

# Lecture 2: Supervised vs. unsupervised learning, bias-variance tradeoff

Reading: Chapter 2


STATS 202: Data mining and analysis

September 25, 2019

## Supervised vs. unsupervised learning

In **unsupervised learning** we seek to understand the relationships between variables in a data matrix:

Samples or observations




Variables or factors

## Supervised vs. unsupervised learning

In **unsupervised learning** we seek to understand the relationships between variables in a data matrix:

Samples or observations



Variables or factors

Quantitative variables, eg. weight, height, number of children, ...

## Supervised vs. unsupervised learning

In **unsupervised learning** we seek to understand the relationships between variables in a data matrix:

Samples or observations



Variables or factors

Qualitative variables, eg. college major, profession, gender, ...

## Supervised vs. unsupervised learning


In **unsupervised learning** we seek to understand the relationships between variables in a data matrix:

Our goal is to:

- ▶ Find meaningful relationships between the variables. **Correlation analysis.**
- ▶ Find low-dimensional representations of the data which make it easy to visualize the variables. **PCA, ICA, isomap, locally linear embeddings, etc.**
- ▶ Find meaningful groupings of the data. **Clustering.**


## Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:



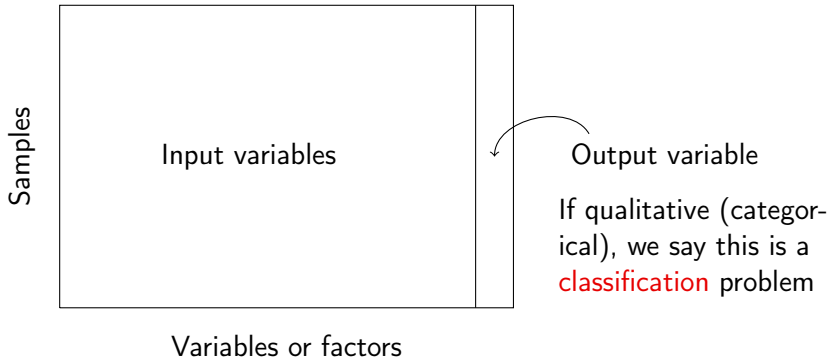
## Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:



## Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:





## Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:

If  $X$  is the vector of inputs for a particular sample. The output variable is modeled by:

$$Y = f(X) + \underbrace{\epsilon}_{\text{Random error}}$$

The function  $f$  captures the systematic relationship between  $X$  and  $Y$ .  $f$  is fixed and unknown.

$\epsilon$  represents the unpredictable "noise" in the problem.

Our goal is to learn the function  $f$ , using a set of **training** samples.

# Supervised vs. unsupervised learning

$$Y = f(X) + \underbrace{\varepsilon}_{\text{Random error}}$$

Motivations:

- **Prediction:** Useful when the input variable is readily available, but the output variable is not.

Example: Predict stock prices next month using data from last year.

- **Inference:** A model for  $f$  can help us understand the structure of the data — which variables influence the output, and which don't? What is the relationship between each variable and the output, e.g. linear, non-linear?

Example: What is the influence of genetic variations on the incidence of heart disease.

## Parametric and nonparametric methods:

Most supervised learning methods fall into one of two classes:



- ▶ **Parametric methods:** We assume that  $f$  takes a specific form. For example, a linear form:

$$f(X) = X_1\beta_1 + \cdots + X_p\beta_p$$

with parameters  $\beta_1, \dots, \beta_p$ . Using the training data, we try to *fit* the parameters.

- ▶ **Non-parametric methods:** We don't make any assumptions on the form of  $f$ , but we restrict how “wiggly” or “rough” the function can be.

## Parametric vs. nonparametric prediction



Figures 2.4 and 2.5

Parametric methods have a limit of fit quality. Non-parametric methods keep improving as we add more data to fit.

Parametric methods are often simpler to interpret.

## Prediction error

**Training data:**  $(x_1, y_1), (x_2, y_2) \dots (x_n, y_n)$

**Predicted function:**  $\hat{f}$ , which is based on these data.

What is a good choice for  $\hat{f}$ ?

Our goal in supervised learning is to minimize the **prediction error**:

For a new datapoint  $(x_0, y_0)$  (not used in training) we want the squared error  $(y_0 - \hat{f}(x_0))^2$  to be small.

Given many test data  $\{(x'_i, y'_i); i = 1, \dots, m\}$  which were not used to fit the model, a common measure of quality of  $\hat{f}$  is the test mean squared error (MSE):

$$MSE_{\text{test}}(\hat{f}) = \frac{1}{m} \sum_{i=1}^m (y'_i - \hat{f}(x'_i))^2.$$


## Prediction error

What if we don't have any test data? It is tempting to use instead the training MSE:

$$MSE_{\text{training}}(\hat{f}) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2.$$


The main challenge of statistical learning is that *a low training MSE does not imply a low test MSE*.

Figure 2.9.



The circles are simulated data from the black curve. In this artificial example, we *know* what  $f$  is.

Figure 2.9.

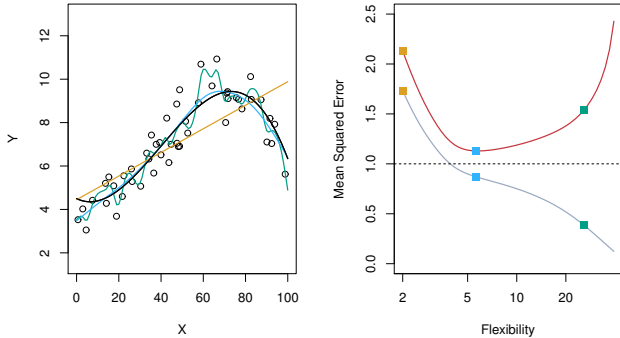


Three estimates  $\hat{f}$  are shown:

1. Linear regression.
2. Splines (very smooth).
3. Splines (quite rough).




Figure 2.9.




Red line: Test MSE.  
Gray line: Training MSE.

Figure 2.10



The function  $f$  is now almost linear.

Figure 2.11



When the noise  $\varepsilon$  has small variance, the third method does well.

## The bias variance decomposition

Let  $x_0$  be a fixed test point,  $y_0 = f(x_0) + \varepsilon_0$ , and  $\hat{f}$  be estimated from  $n$  training samples  $(x_1, y_1) \dots (x_n, y_n)$ .

Let  $E$  denote the expectation over  $y_0$  and the training data. Then, the expected test MSE at  $x_0$  can be decomposed:

$$E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$

## The bias variance decomposition

Let  $x_0$  be a fixed test point,  $y_0 = f(x_0) + \varepsilon_0$ , and  $\hat{f}$  be estimated from  $n$  training samples  $(x_1, y_1) \dots (x_n, y_n)$ .

Let  $E$  denote the expectation over  $y_0$  and the training data. Then, the expected test MSE at  $x_0$  can be decomposed:

$$E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$

Irreducible error

## The bias variance decomposition

Let  $x_0$  be a fixed test point,  $y_0 = f(x_0) + \varepsilon_0$ , and  $\hat{f}$  be estimated from  $n$  training samples  $(x_1, y_1) \dots (x_n, y_n)$ .

Let  $E$  denote the expectation over  $y_0$  and the training data. Then, the expected test MSE at  $x_0$  can be decomposed:

$$E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$

The variance of the estimate of  $Y$ :  $E[\hat{f}(x_0) - E(\hat{f}(x_0))]^2$

This measures how much the estimate of  $\hat{f}$  at  $x_0$  changes when we sample new training data.

## The bias variance decomposition





Let  $x_0$  be a fixed test point,  $y_0 = f(x_0) + \varepsilon_0$ , and  $\hat{f}$  be estimated from  $n$  training samples  $(x_1, y_1) \dots (x_n, y_n)$ .

Let  $E$  denote the expectation over  $y_0$  and the training data. Then, the expected test MSE at  $x_0$  can be decomposed:





$$E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$







The squared bias of the estimate of  $Y$ :  $[E(\hat{f}(x_0)) - f(x_0)]^2$

This measures the deviation of the average prediction  $\hat{f}(x_0)$  from the truth  $f(x_0)$ .









## Implications of bias variance decomposition

$$E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon).$$


- ▶ The expected MSE is never smaller than the irreducible error.
- ▶ Both small bias and small variance are needed for small expected MSE.
- ▶ Easy to find zero variance procedure with high bias (predict a constant value) and very low bias procedure with high variance (choose  $\hat{f}$  passing through every training point).
- ▶ In practice, best expected MSE achieved by incurring some bias to decrease variance and vice-versa: this is the **bias-variance trade-off**.

### Rule of thumb:


More flexible methods  $\Rightarrow$  **Higher variance** and **Lower Bias**.

- ▶ **Example:** Linear fit may not change substantially with training data but has high bias if  $f$  is very non-linear.

Squiggly  $f$ , high noise



Linear  $f$ , high noise



Squiggly  $f$ , low noise




Figure 2.12

# Classification problems

In a classification setting, the output takes values in a discrete set.

For example, if we are predicting the brand of a car based on a number of variables, the function  $f$  takes values in the set  $\{\text{Ford, Toyota, Mercedes-Benz, } \dots\}$ .

We adopt some additional notation:

- $P(X, Y)$  : joint distribution of  $(X, Y)$ ,
- $P(Y \mid X)$  : conditional distribution of  $Y$  given  $X$ ,
- $\hat{y}_i$  : prediction for  $x_i$ .

## Loss function for classification

There are many ways to measure the error of a classification prediction. One of the most common is the 0-1 loss:

$$\mathbf{1}(y_0 \neq \hat{y}_0)$$

As with squared error, we can compute average test prediction error (called **test error rate** under 0-1 loss) using previously unseen test data  $\{(x'_i, y'_i); i = 1, \dots, m\}$ :

$$\frac{1}{m} \sum_{i=1}^m \mathbf{1}(y'_i \neq \hat{y}'_i)$$

Similarly, we can compute the (usually optimistic) **training error rate**

$$\frac{1}{n} \sum_{i=1}^n \mathbf{1}(y_i \neq \hat{y}_i)$$

## Bayes classifier




Figure 2.13

In practice, we never know the joint probability  $P$ . However, we can assume that it exists.

The **Bayes classifier** assigns:

$$\hat{y}_i = \operatorname{argmax}_j P(Y = j \mid X = x_i)$$

It can be shown that this is the best classifier under the 0-1 loss.